

3DV-2023 Assignment: Single View to 3D

3DV-2023 Assignment: Single View to 3D

Goal

Detailed explanation of your code

Dataset

Loss

Model

SingleViewto3D

PointDecoder

VoxelDecoder

Visualization

Train

Eval

Visualization and analysis of your experiment result

Point cloud

Settings

ResNet18

Resnet34

ResNet50

Voxel

Settings

ResNet18

ResNet50

⚡ PDF provides index navigation. You can utilize it to navigate to the part you want to read.

Goal

In this assignment, you will explore the types of loss and decoder functions for regressing to voxels, and point clouds representation from single view RGB input.

Detailed explanation of your code

Dataset

Read images and point clouds/voxels from files.

The code I changed:

1. Remove unused code comments
2. Remove permuting image shape from (B, C, H, W) to (B, H, W, C)

In torch tensor's common format, it is (B, C, H, W).

In this training flow, I find only the model use the images, and it also permutes back to (B, C, H, W).

So, It is unnecessary.

Point cloud format: (B, N, 3)

N: number of points in one image

3: xyz coordinate in 3D space.

Voxel format: binary format, (B, A, B, C)

A, B, C: the size of voxels (D, H, W)

```
class ShapeNetDB(Dataset):
    def __init__(self, data_dir: str, data_type: str, n_points: int = 2048):
        super().__init__()
        self.data_dir = data_dir
        self.data_type = data_type
        self.n_points = n_points
        self.db = self.load_db()

        self.get_index()

    def __len__(self):
        return len(self.db)

    def __getitem__(self, idx):
        if self.data_type == "point":
            """
            Return shapes:
            img: (B, 256, 256, 3)
            pc: (B, 2048, 3)
            object_id: (B,)
            """
            img, img_id = self.load_img(idx)
            pc, object_id = self.load_point(idx)

            assert img_id == object_id

            return img, pc, object_id

        elif self.data_type == "voxel":
            """
            Return shapes:
            img: (B, 256, 256, 3)
            voxel: (B, 33, 33, 33)
            object_id: (B,)
            """
            img, img_id = self.load_img(idx)
            voxel, object_id = self.load_voxel(idx)

            assert img_id == object_id

            return img, voxel, object_id

    def load_db(self):
        db_list = sorted(glob.glob(os.path.join(self.data_dir, "*")))

        return db_list

    def get_index(self):
        self.id_index = self.data_dir.split("/").index("data") + 2

    def load_img(self, idx):
        path = os.path.join(self.db[idx], "view.png")
```

```

img = read_image(path) / 255.0

object_id = self.db[idx].split("/")[self.id_index]

return img, object_id

def load_point(self, idx):
    path = os.path.join(self.db[idx], "point_cloud.npy")
    points: np.ndarray = np.load(path)

    # resample
    # if self.n_points < points.shape[0]:
    #     choice = np.random.choice(points.shape[0], self.n_points,
replace=False)
    #     points = points[choice]

    # normalize
    points -= np.mean(points, axis=0, keepdims=True) # center
    dist = np.max(np.sqrt(np.sum(points**2, axis=1)), 0)
    points = points / dist # scale

    object_id = self.db[idx].split("/")[self.id_index]

    return torch.from_numpy(points), object_id

def load_voxel(self, idx):
    path = os.path.join(self.db[idx], "voxel.npy")
    voxel = np.load(path)

    object_id = self.db[idx].split("/")[self.id_index]

    return torch.from_numpy(voxel).float(), object_id

```

Loss

voxel_loss: Binary cross entropy with logits (contain sigmoid function)

chamfer_loss: Use the nearest points to calculate loss.

$$d_{CD}(X, Y) = \sum_{x \in X} \min_{y \in Y} \|x - y\|_2^2 + \sum_{y \in Y} \min_{x \in X} \|x - y\|_2^2$$

```

def voxel_loss(voxel_src: torch.Tensor, voxel_tgt: torch.Tensor):
    loss = binary_cross_entropy_with_logits(voxel_src, voxel_tgt)
    return loss

def chamfer_loss(point_cloud_src: torch.Tensor, point_cloud_tgt: torch.Tensor):
    assert point_cloud_src.ndimension() == 3 # (B, N, 3)
    assert point_cloud_src.size(-1) == 3
    assert point_cloud_tgt.ndimension() == 3 # (B, N, 3)
    assert point_cloud_tgt.size(-1) == 3

    # [B, N, N, 3]
    distance = point_cloud_src[:, :, None, :] - point_cloud_tgt[:, None, :, :]
    # [B, N, N]

```

```

distance = torch.sum(distance**2, dim=-1)

# [B, N]
min_xy, _ = torch.min(distance, dim=-1)
min_yx, _ = torch.min(distance.transpose(1, 2), dim=-1)

# [B]
loss_xy = min_xy.mean(dim=1)
loss_yx = min_yx.mean(dim=1)

return (loss_xy + loss_yx).mean()

```

Model

SingleViewto3D

An Encoder-Decoder architecture.

Encoder uses ResNet series.

The code I changed:

1. Freeze the BatchNorm2d layers

While using ResNet50 as encoder, I found the batch normalization will make the performance bad.

Batch Normalization will be determined by the batch size.

If the batch size is very small (e.g. 16, 8...), it will perform not well.

2. Build the model automatically.

I find it is determined by BasicBlock or Bottleneck.

Refer to [here](#).

```

class SingleViewto3D(nn.Module):
    def __init__(self, cfg):
        super().__init__()
        self.dtype = cfg.dtype
        vision_model: torch.nn.Module = torchvision_resnet.__dict__[cfg.arch](
            pretrained=True
        )
        self.encoder = torch.nn.Sequential(*(list(vision_model.children())[:-1]))
        for module in self.encoder.modules():
            if isinstance(module, nn.BatchNorm2d):
                module.requires_grad_(False)

        latent_size = 512
        for module in vision_model.modules():
            if isinstance(module, torchvision_resnet.BasicBlock):
                latent_size *= torchvision_resnet.BasicBlock.expansion
                break
            elif isinstance(module, torchvision_resnet.Bottleneck):
                latent_size *= torchvision_resnet.Bottleneck.expansion
                break

        # define decoder
        if cfg.dtype == "voxel":
            self.decoder = VoxelDecoder(cfg.n_points, latent_size)

```

```

elif cfg.dtype == "point":
    self.decoder = PointDecoder(cfg.n_points, latent_size)

self.normalize = Normalize(
    mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225]
)

def forward(self, images: torch.Tensor):
    images = self.normalize(images)
    encoded_feat = self.encoder(images)
    encoded_feat = torch.flatten(encoded_feat, 1)
    pred = logits = self.decoder(encoded_feat)
    if self.dtype == "voxel":
        pred = torch.sigmoid(logits)
    return logits, pred

```

PointDecoder

The code I changed:

1. Build the layers automatically.
2. Add Group Normalization to let the model coverage more easily.

```

class PointDecoder(nn.Module):
    def __init__(self, num_points: int, latent_size: int = 512):
        super().__init__()
        self.num_points = num_points
        out_features = self.num_points * 3

        layers = []
        scale = 2
        max_scale = out_features // latent_size
        prev_size = latent_size
        while scale < max_scale:
            next_size = latent_size * scale
            layers.extend(
                [
                    nn.Linear(prev_size, next_size),
                    nn.GroupNorm(32, num_channels=next_size),
                ]
            )
            prev_size = next_size
            scale *= 2

        layers.append(nn.Linear(prev_size, out_features))

        self.layers = nn.ModuleList(layers)
        self.th = nn.Tanh()

    def forward(self, x: torch.Tensor):
        for layer in self.layers[:-1]:
            x = F.relu(layer(x))

        x = self.layers[-1](x)
        x = self.th(x)

```

```
x = x.view(-1, self.num_points, 3)
return x
```

VoxelDecoder

The code I changed:

1. Remove Tanh activation layer.

Because of binary format of voxels, the decoder predicts multiple binary score.

Tanh is used for predicting (-1, 1) values, such as coordinates.

So, I change to use sigmoid with binary entropy loss function.

2. Add Group Normalization to let the model coverage more easily.

```
class VoxelDecoder(nn.Module):
    def __init__(self, num_points: int, latent_size: int = 512):
        super().__init__()
        self.num_points = num_points
        self.fc0 = nn.Linear(latent_size, latent_size * 2)
        self.gn0 = nn.GroupNorm(32, num_channels=latent_size * 2)
        self.fc1 = nn.Linear(latent_size * 2, latent_size * 4)
        self.gn1 = nn.GroupNorm(32, num_channels=latent_size * 4)

        self.fc5 = nn.Linear(latent_size * 4, self.num_points**3)

    def forward(self, x: torch.Tensor):
        x = self.fc0(x)
        x = self.gn0(x)
        x = F.relu(x)
        x = self.fc1(x)
        x = self.gn1(x)
        x = F.relu(x)

        x = self.fc5(x)
        x = x.view(-1, self.num_points, self.num_points, self.num_points)
        return x
```

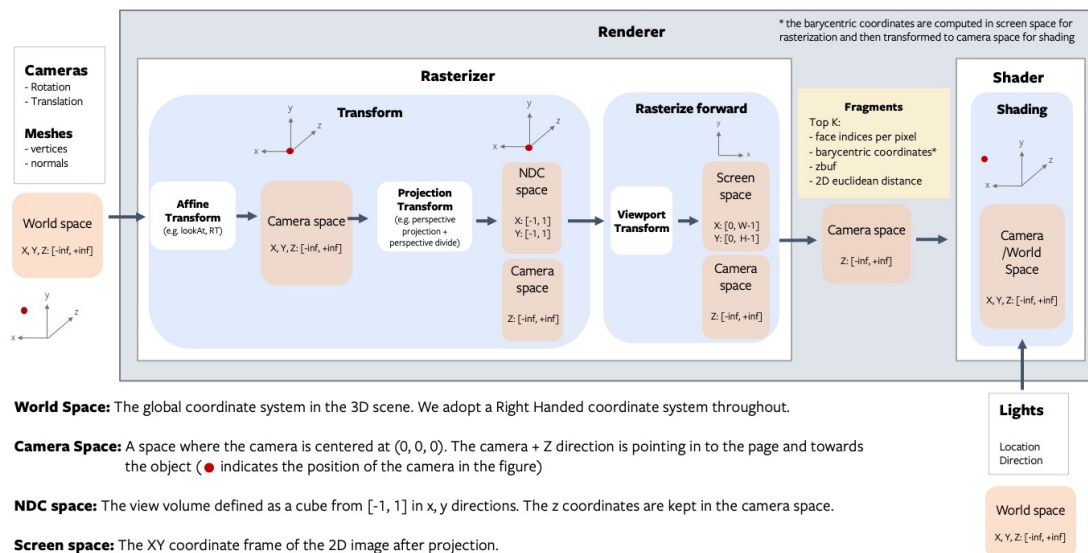
Visualization

Refer to the script scene.py in notebooks folder.

Use pytorch3d.renderer to visualize the point clouds and voxels.

Later, in eval part, I will explain how to convert binary format of voxels to cubes.

Rendering flow ([renderer getting started · PyTorch3D](#))



```
class Voxelscene:
    def __init__(self, device: torch.device):
        self.device = device

    def set_cam(self, dist=1.0, elev=0.0, azim=0.0):
        """
        Initialize a camera.
        with world coordinates +Y up, +X left and +Z in.
        see full api in
        https://pytorch3d.readthedocs.io/en/latest/modules/renderer/cameras.html#pytorch3d.renderer.cameras.look_at_view_transform
        """
        R, T = look_at_view_transform(dist, elev, azim)
        self.cameras = FoVPerspectiveCameras(device=self.device, R=R, T=T)
        return self

    def set_light(self, location=[[0.0, 0.0, 0.0]]):
        """
        see full api in
        https://pytorch3d.readthedocs.io/en/latest/modules/renderer/lighting.html#pytorch3d.renderer.lighting.PointLights
        """
        self.lights = PointLights(location=location, device=self.device)
        return self

    def set_rasterizer(self, image_size=512):
        self.raster_settings = RasterizationSettings(
            image_size=image_size,
            blur_radius=1e-5,
            faces_per_pixel=150,
            bin_size=0,
            cull_backfaces=True,
            # max_faces_per_bin=10,
        )
```

```

        see full api in
https://pytorch3d.readthedocs.io/en/latest/modules/renderer/rasterizer.html#pytorch3d.renderer.mesh.rasterizer.RasterizationSettings
        """

        return self

    def set_renderer(self):
        """
        see full api in
https://github.com/facebookresearch/pytorch3d/blob/2c64635daa2aa728f35ed4abe41c6942ae8c0d8b/pytorch3d/renderer/mesh/renderer.py#L32
        """

        self.renderer = MeshRenderer(
            rasterizer=MeshRasterizer(
                cameras=self.cameras, raster_settings=self.raster_settings
            ),
            shader=SoftGouraudShader(
                device=self.device, cameras=self.cameras, lights=self.lights
            ),
        )

        return self

class PointScene:
    def __init__(self, device: torch.device):
        self.device = device

    def set_cam(self, dist=1.0, elev=0.0, azimuth=0.0):
        """
        Initialize a camera.
        with world coordinates +Y up, +X left and +Z in.
        see full api in
https://pytorch3d.readthedocs.io/en/latest/modules/renderer/cameras.html#pytorch3d.renderer.cameras.look\_at\_view\_transform
        """
        R, T = look_at_view_transform(dist, elev, azimuth)
        self.cameras = FoVPerspectiveCameras(device=self.device, R=R, T=T)
        return self

    def set_rasterizer(self, image_size=512):
        self.raster_settings = PointsRasterizationSettings(
            image_size=image_size,
            radius=1e-2,
            points_per_pixel=120,
            bin_size=0,
            # max_points_per_bin=10,
        )
        """
        see full api in
https://pytorch3d.readthedocs.io/en/latest/modules/renderer/rasterizer.html#pytorch3d.renderer.mesh.rasterizer.RasterizationSettings
        """
        return self

```



```

def set_renderer(self):
    """
    see full api in
    https://github.com/facebookresearch/pytorch3d/blob/2c64635daa2aa728f35ed4abe41c69
    42ae8c0d8b/pytorch3d/renderer/mesh/renderer.py#L32
    """

    self.renderer = PointsRenderer(
        rasterizer=PointsRasterizer(
            cameras=self.cameras, raster_settings=self.raster_settings
        ),
        compositor=NormWeightedCompositor(background_color=[1.0, 1.0, 1.0]),
    )

    return self

```

Train

The code I changed:

1. Change Adam to AdamW
2. Visualize loss history
3. Choose GPU device from config
4. Shuffle the data for each epoch

```

def calculate_loss(predictions, ground_truth, cfg):
    if cfg.dtype == "voxel":
        loss = losses.voxel_loss(predictions, ground_truth)
    elif cfg.dtype == "point":
        loss = losses.chamfer_loss(predictions, ground_truth)
    return loss

@hydra.main(config_path="configs/", config_name="config.yaml")
def train_model(cfg: DictConfig):
    log.info(f"Device: {cfg.device}")
    device = torch.device(cfg.device)

    log.info(cfg.data_dir)
    shapenetdb = ShapeNetDB(cfg.data_dir, cfg.dtype)

    loader = torch.utils.data.DataLoader(
        shapenetdb,
        batch_size=cfg.batch_size,
        num_workers=cfg.num_workers,
        pin_memory=True,
        drop_last=False,
        shuffle=True,
    )
    train_loader = iter(loader)

    if cfg.dtype == "voxel":
        cfg.n_points = shapenetdb[0][1].shape[1]

```

```

model = Singleviewto3D(cfg)
model.cuda(device)
model.train()

# ===== preparing optimizer ... =====
optimizer = torch.optim.AdamW(
    model.parameters(), lr=cfg.lr, weight_decay=cfg.weight_decay
) # to use with ViTs
start_iter = 0
start_time = time.time()

if cfg.load_checkpoint:
    checkpoint = torch.load(f"{cfg.base_dir}/checkpoint_{cfg.dtype}.pth")
    model.load_state_dict(checkpoint["model_state_dict"])
    optimizer.load_state_dict(checkpoint["optimizer_state_dict"])
    start_iter = checkpoint["step"]
    log.info(f"Successfully loaded iter {start_iter}")

log.info("Starting training !")
best_loss = sys.float_info.max
loss_history = []
for step in range(start_iter, cfg.max_iter):
    iter_start_time = time.time()

    if step % len(train_loader) == 0: # restart after one epoch
        train_loader = iter(loader)

    read_start_time = time.time()

    images_gt, ground_truth_3d, _ = next(train_loader)
    images_gt, ground_truth_3d = images_gt.cuda(device),
ground_truth_3d.cuda(
    device
)

    read_time = time.time() - read_start_time

    prediction_logits: torch.Tensor
    prediction_logits, _ = model(images_gt)

    loss = calculate_loss(prediction_logits, ground_truth_3d, cfg)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    total_time = time.time() - start_time
    iter_time = time.time() - iter_start_time

    loss_vis = loss.detach().cpu().item()

    if loss_vis < best_loss:
        best_loss = loss_vis
        torch.save(
            {
                "step": step,

```

```

        "model_state_dict": model.state_dict(),
        "optimizer_state_dict": optimizer.state_dict(),
    },
    f"{cfg.base_dir}/checkpoint_{cfg.dtype}.pth",
)

# if (step % cfg.save_freq) == 0:
#     torch.save(
#         {
#             "step": step,
#             "model_state_dict": model.state_dict(),
#             "optimizer_state_dict": optimizer.state_dict(),
#         },
#         f"{cfg.base_dir}/checkpoint_{cfg.dtype}.pth",
#     )

log.info(
    "[%4d/%4d]; ttime: %.0f (%.2f, %.2f); loss: %.5f"
    % (step, cfg.max_iter, total_time, read_time, iter_time, loss_vis)
)

loss_history.append(loss_vis)

log.info("Done!")

log.info("Best loss: %.5f" % (best_loss))

fig = plt.figure()
axe = fig.add_subplot()
axe.plot(loss_history)
axe.set_title("Training Loss")
axe.set_xlabel("Iteration")
axe.set_ylabel("Loss")
fig.savefig(f"loss_{cfg.dtype}.jpg")
plt.close(fig)

```

Eval

The code I changed:


1. Choose GPU device from config
2. Implement visualization functions

Point cloud visualization (`plot_points`) function:

1. Convert points to `pytorch3d.structures.Pointclouds`
2. Use `PointsRenderer` to visualize gt and prediction.

Voxel visualization (`plot_voxels`) function):

1. Cubify the voxels into Meshes ([cubify](#) function is provided by `pytorch3d.ops`)
I use center align.



2. Add color textures to meshes

3. Use MeshRenderer to visualize gt and prediction.

```
def calculate_loss(predictions, ground_truth, cfg):
    if cfg.dtype == "voxel":
        loss = losses.voxel_loss(predictions, ground_truth)
    elif cfg.dtype == "point":
        loss = losses.chamfer_loss(predictions, ground_truth)
    return loss

def plot_points(predictions: torch.Tensor) -> torch.Tensor:
    normals = estimate_pointcloud_normals(predictions)
    point_clouds = Pointclouds(
        points=predictions,
        normals=normals,
        features=torch.full_like(predictions, 0.5, device=predictions.device),
    )

    scene = PointScene(predictions.device)
    scene.set_cam(1.8, 45.0, 45.0)
    scene.set_rasterizer(image_size=256)
    scene.set_renderer()
    images = scene.renderer(point_clouds)
    return images

def plot_voxels(predictions: torch.Tensor):
    meshes = cubify(predictions, thresh=0.5, align="center")

    verts_list = meshes.verts_packed()
    verts_rgb = torch.ones(1, verts_list.shape[0], 3, device=predictions.device)
    meshes.textures = TexturesVertex(verts_features=verts_rgb)

    scene = VoxelScene(predictions.device)
    scene.set_cam(2.5, 45.0, 45.0)
    scene.set_light(location=[[0.0, 0.0, 0.0]])
    scene.set_rasterizer(image_size=256)
    scene.set_renderer()
    images = scene.renderer(meshes)
    return images
```

@torch.no_grad()

```

@hydra.main(config_path="configs/", config_name="config.yaml")
def evaluate_model(cfg: DictConfig):
    device = torch.device(cfg.device)

    shapenetdb = ShapeNetDB(cfg.data_dir, cfg.dtype, cfg.n_points)

    loader = torch.utils.data.DataLoader(
        shapenetdb,
        batch_size=cfg.batch_size,
        num_workers=cfg.num_workers,
        pin_memory=True,
        drop_last=True,
        shuffle=False,
    )
    eval_loader = iter(loader)

    if cfg.dtype == "voxel":
        cfg.n_points = shapenetdb[0][1].shape[1]

    model = Singleviewto3D(cfg)
    model.cuda(device)
    model.eval()

    start_iter = 0
    start_time = time.time()

    avg_loss = []

    if cfg.load_eval_checkpoint:
        checkpoint = torch.load(f"{cfg.base_dir}/checkpoint_{cfg.dtype}.pth")
        model.load_state_dict(checkpoint["model_state_dict"])
        log.info(f"Successfully loaded iter {start_iter}")

    vis_dir = os.path.join(cfg.base_dir, "vis")
    os.makedirs(vis_dir, exist_ok=True)

    log.info("Starting evaluating !")
    max_iter = len(eval_loader)
    for step in range(start_iter, max_iter):
        iter_start_time = time.time()

        read_start_time = time.time()

        images_gt, ground_truth_3d, _ = next(eval_loader)
        images_gt, ground_truth_3d = images_gt.cuda(device),
        ground_truth_3d.cuda(
            device
        )

        read_time = time.time() - read_start_time

        prediction_logits: torch.Tensor
        prediction_3d: torch.Tensor
        prediction_logits, prediction_3d = model(images_gt)
        torch.save(prediction_3d.cpu(), f"{cfg.base_dir}/pre_point_cloud.pt")

```

```

loss = calculate_loss(prediction_logits, ground_truth_3d,
cfg).cpu().item()

if (step % cfg.vis_freq) == 0:
    # visualization block
    if cfg.dtype == "point":
        gt_images = plot_points(ground_truth_3d)
        images = plot_points(prediction_3d)
    elif cfg.dtype == "voxel":
        gt_images = plot_voxels(ground_truth_3d)
        images = plot_voxels(prediction_3d)

    rgb_images = images_gt.permute(0, 2, 3, 1).cpu().numpy()
    gt_images = gt_images.cpu().numpy()
    images = images.cpu().numpy()

    f = plt.figure()
    for i in range(images.shape[0]):
        rgb_image = rgb_images[i]
        gt_image = gt_images[i]
        image = images[i]

        axeses = f.subplots(1, 3, sharey=True)
        for axes in axeses:
            axes.xaxis.set_visible(False)
            axes.yaxis.set_visible(False)

        ax1, ax2, ax3 = axeses
        ax1.set_title("GT Image")
        ax1.imshow(rgb_image)

        ax2.set_title("GT 3D")
        ax2.imshow(gt_image)

        ax3.set_title("Prediction 3D")
        ax3.imshow(image)

        file_name = os.path.join(vis_dir, f"{step}_{cfg.dtype}_{i}.png")
        f.savefig(file_name)
        f.clf()
    plt.close(f)

total_time = time.time() - start_time
iter_time = time.time() - iter_start_time

avg_loss.append(loss)

log.info(
    "[%4d/%4d]; ttime: %.0f (%.2f, %.2f); eva_loss: %.3f"
    % (
        step,
        max_iter,
        total_time,
        read_time,
        iter_time,
        torch.tensor(avg_loss).mean(),
    )

```

```
log.info("Done!")
```

Visualization and analysis of your experiment result

Point cloud

Settings

- learning rate: 0.001
- weight decay: 0.01
- max iterations: 20000
- batch size: 32(ResNet18, Resnet34), 16(ResNet50, ResNet101)
- num workers: 4

ResNet18


Best Loss: 0.00271

Although the loss is very low, the points spread every where.

This also explains why chamfer loss is not the best option.

Chamfer loss is using the nearest point to calculate. It does not filter out the selected points.

So, some points may be constrained by loss function many times.



GT Image



GT 3D




Prediction 3D




GT Image



GT 3D



Prediction 3D




GT Image




GT 3D



Prediction 3D






Resnet34

Best Loss: 0.00253

Same problem.


Although the loss is very low, the points spread every where.



GT Image



GT 3D



Prediction 3D




GT Image



GT 3D



Prediction 3D




GT Image



GT 3D



Prediction 3D






ResNet50

Best Loss: 0.00276

Same problem.

Although the loss is very low, the points spread every where.


From these results, I think it coverage.




GT Image



GT 3D



Prediction 3D




GT Image



GT 3D



Prediction 3D



GT Image



GT 3D



Prediction 3D





Voxel


Settings

- learning rate: 0.001
- weight decay: 0.01
- max iterations: 20000
- batch size: 256(ResNet18), 32(ResNet50)
- num workers: 4

ResNet18

Best Loss: 0.00007


The model overfits(maybe) perfectly.




GT Image



GT 3D



Prediction 3D



GT Image



GT 3D




Prediction 3D




GT Image




GT 3D



Prediction 3D






ResNet50

Best Loss: 0.01257

Due to not enough GPU memory space, the batch size is decreased.

So, the model needs more iterations to coverage.


But I think it will be like previous one overfitting perfectly 🙌 .




GT Image



GT 3D




Prediction 3D



GT Image



GT 3D




Prediction 3D




GT Image



GT 3D




Prediction 3D



GT Image



GT 3D



Prediction 3D

